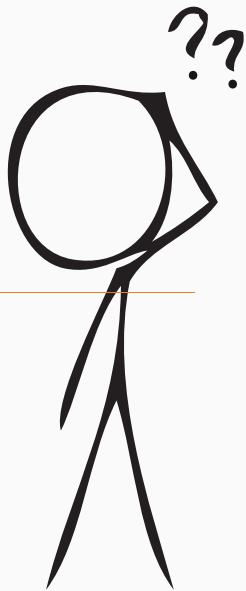


VERIFICATION OF CONCURRENT AND DISTRIBUTED SYSTEMS

Nickolai Novik

April 18, 2018

<http://github.com/jettify>



ABOUT ME

- **[Software Engineer]** DataRobot
- **[Github]** <http://github.com/jettify>
- **[Twitter]** <https://twitter.com/isinf>
- **[aio-lib]** <https://github.com/aio-lib>
- **[Projects]** *aiomonitor, aiohttp-debugtoolbar, aiobotocore, aiomysql, aiiodbc, aiohttp-admin, aiowlock, aiozipkin, etc*

AGENDA

1. Problem Statement. Motivational Example
2. Model Checking
3. TLA+ Basics
4. Multithreading Queue Spec
5. Aiorwlock Spec
6. Conclusions

How many of you heard of formal methods?

- I used one on of: Coq, Isable, TLA+, Alloy, Spin.
- I heard about it, but never used.
- I think formal methods are kinda cool.

PROBLEM STATEMENT. MOTIVATIONAL EXAMPLE

CONCURRENT AND DISTRIBUTED SYSTEMS ARE HARD

Most of industry uses following techniques for quality assurance:

1. Design review
2. Static code analysis
3. Unit/Integration/Functional testing
4. Code coverage
5. Code review
6. Stress testing
7. Fault-injection testing [1]

DISTRIBUTED ALGORITHMS EXTREMELY HARD

Chord: A Scalable Peer-to-peer Lookup Protocol for Internet Applications

Ion Stoica[†], Robert Morris[‡], David Liben-Nowell[‡], David R. Karger[‡], M. Frans Kaashoek[‡], Frank Dabek[‡], Hari Balakrishnan[‡]

Abstract—

A fundamental problem that confronts peer-to-peer applications is the efficient location of the node that stores a desired data item. This paper presents *Chord*, a distributed lookup protocol that addresses this problem. *Chord* provides support for just one operation: given a key, it maps the key onto a node. Data location can be easily implemented on top of *Chord* by associating a key with each data item, and storing the key/data pair at the node to which the key maps. *Chord* adapts efficiently as nodes join and leave the system, and can answer queries even if the system is continuously changing. Results from theoretical analysis and simulations show that *Chord* is scalable: communication cost and the state maintained by each node scale logarithmically with the number of *Chord* nodes.


tem.

A *Chord* node requires information about $O(\log N)$ other nodes for *efficient* routing, but performance degrades gracefully when that information is out of date. This is important in practice because nodes will join and leave arbitrarily, and consistency of even $O(\log N)$ state may be hard to maintain. Only one piece of information per node need be correct in order for *Chord* to guarantee correct (though possibly slow) routing of queries; *Chord* has a simple algorithm for maintaining this information in a dynamic environment.

Chord

is popular algorithm for P2P systems, paper published in 2001 by strong team of MIT researchers, 10 years later bug found in specification [11, 12]. Paper won best paper award.

SOMETIMES COST OF ERROR IS VERY HIGH



Mars Pathfinder

rover, the mission was jeopardised by a concurrent software bug in the lander. [10]

SOMETIMES COST OF ERROR IS VERY HIGH 2




Therac-25

radiation therapy machine, because of concurrent programming errors, it sometimes gave its patients radiation doses that were hundreds of times greater than normal [2]

MODEL CHECKING

FORMAL METHODS

Formal Methods - techniques and tools based on mathematics and formal logic [8].




MODEL CHECKING

Model checking is a technique for automatically verifying correctness properties of finite-state systems.


model checking:

all program state are explored
until none left or defect found




testing:

based on input set $\{d\}$
only one **path**
executed at a time



TLA+ – TEMPORAL LOGIC OF ACTIONS



TLA+ language developed by **Leslie Lamport**. It is used to design, model, document, and verify concurrent systems, has been described as exhaustively-testable pseudocode and blueprints for software systems.

TLA+ DESCRIPTION

TLA+ is **specifications** language, it is precise written description of what a system suppose to do. **TLA Toolbox** distribution of TLA related tools. **TLC** - model checker tool to validate invariants stated in spec. **TLAPS** - mechanical proof checker.

Released April 23, 1999; 18 years ago

Syntax math notation, similar to \LaTeX

Implem. Java

License MIT


IDE TLA Toolbox (Eclipse based) [5]



TLA+ helped to find design bugs in **S3**, **Dynamo**, **EBS**, **EC2**, etc, some requiring traces of 35 steps. [9]

TLA+ INDUSTRY USAGE: MICROSOFT

MS used TLA+ to define consistency protocol for CosmosDB and memory allocator for Xbox [7]



TLA+ INDUSTRY USAGE: OPEN-SOURCE

Open source projects use TLA+ to verify algorithms:

- **Linux Kernel** – verify fairness of qspinlock [6]
- **Elastic** – data replication protocol [3]
- **Mongodb** – data replication protocol [13]
- **Hadoop/YARN** – registry of long lived processes [4]



RAFT AND TLA+

```
452  Next == /\ \/ \E i \in Server : Restart(i)
453          \/ \E i \in Server : Timeout(i)
454          \/ \E i,j \in Server : RequestVote(i, j)
455          \/ \E i \in Server : BecomeLeader(i)
456          \/ \E i \in Server, v \in Value : ClientRequest(i, v)
457          \/ \E i \in Server : AdvanceCommitIndex(i)
458          \/ \E i,j \in Server : AppendEntries(i, j)
459          \/ \E m \in DOMAIN messages : Receive(m)
460          \/ \E m \in DOMAIN messages : DuplicateMessage(m)
461          \/ \E m \in DOMAIN messages : DropMessage(m)
```

RAFT - core algorithm of **consul**, **etcd** also verified with TLA+.

TLA+ BASICS

TLA+ HELLO WORLDS

```
1  ----- MODULE HourClock -----
2  EXTENDS Naturals
3  VARIABLE hr
4  HCini == hr \in (1 .. 12)
5  HCnxt == hr' = IF hr # 12 THEN hr + 1 ELSE 1
6  HC == HCini /\ [] [HCnxt]_hr
7  -----
8  THEOREM HC => []HCini
9  =====
```

```
1  |----- MODULE HourClock -----|
2  | EXTENDS Naturals |
3  | VARIABLE hr      |
4  | HCini   $\triangleq$  hr  $\in$  (1 .. 12) |
5  | HCnxt   $\triangleq$  hr' = IF hr  $\neq$  12 THEN hr + 1 ELSE 1 |
6  | HC      $\triangleq$  HCini  $\wedge \square$ [HCnxt]hr |
7  |-----|
8  | THEOREM HC  $\Rightarrow \square$ HCini |
9  |-----|
```

TLA+ SYNTAX. LOGIC

Basic logical operators:

Symb.	Python	Description
\vee	<code>or</code>	logical OR
\wedge	<code>and</code>	logical AND
\neg	<code>not</code>	logical NOT
$=$	<code>==</code>	boolean operator, checks equality, it is not an assignment operator.
\triangleq	<code>=</code>	means defined to equal.

EXERCISE

$$FALSE \wedge (TRUE \wedge (FALSE \vee TRUE) \vee (TRUE \vee FALSE)) = \boxed{??}$$

EXERCISE

$$\boxed{FALSE} \wedge (TRUE \wedge (FALSE \vee TRUE) \vee (TRUE \vee FALSE)) = \boxed{FALSE}$$

TLA+ SYNTAX. MORE LOGIC

More logic operators:

Symb.	Python	Description
\exists	<code>any()</code>	means "there exists", written as \exists in ASCII
\forall	<code>all()</code>	means "for all", written as \forall in ASCII
:		reads as "such that"

$\exists x \in 1, 2, 3, 4, 5 : x > 3$ - exists x in set of integers 1, 2, 3, 4, 5 such that $x > 3$ expression evaluates to TRUE.

TLA+ SYNTAX. MORE LOGIC

More logic operators:

Operator	Description
\square	formula is TRUE on each step $[]$ in ASCII
\implies	logical implication $x_imp_y = y$ if x else True
$'$	reads as prime, state of variable on next step $x = x + 1$

$Init \wedge [] [Next]_{hr}$ formula true on each step for temporal variable hr

TLA+ SYNTAX. SETS

Basic set operations:

Symb.	Python	Description
$S \cup T$	<code>s.union(t)</code>	Union
$S \cap T$	<code>s.intersection(t)</code>	Intersection
$S \subseteq T$	<code>s in t</code>	Membership
$S \setminus T$	<code>s.difference(t)</code>	Difference


TLA+ SPEC TEMPLATE

```
1  ----- MODULE ModuleName -----
2  (* Imports and variable declarations*)
3  EXTENDS Naturals, Sequences, Integers, FiniteSets
4  -----
5  (* Initial Conditions *)
6  Init == ...
7  TypeOK == ...
8  -----
9  (* Body of the spec *)
10 Next == ...
11 Invariant == ...
12 -----
13 (* Invariant declaration with temporal formula *)
14 THEOREM Spec => ...
15 =====
```

MULTITHREADING QUEUE SPEC

MULTITHREADING QUEUE

Classic *bounded buffer*, attempts to put an element into a full queue or take from empty will block.



QUEUE IMPLEMENTATION, PART 1

```
1  import threading
2
3  class BoundedQueue:
4      def __init__(self, capacity=3):
5          self.capacity = capacity
6          self.buffer = [None] * capacity
7          self.mutex = threading.Lock()
8          self.condition = threading.Condition(self.mutex)
9          self.size = 0
10         self.head = 0
11         self.tail = 0
12
13     def is_full(self):
14         return self.size == self.capacity
15
16     def is_empty(self):
17         return self.size == 0
18
19     def _next(self, x):
20         return (x + 1) % self.capacity
```

QUEUE IMPLEMENTATION, PART 2

```
22     def put(self, item):
23         with self.condition:
24             while self.is_full():
25                 self.condition.wait()
26             self.buffer[self.tail] = item
27             self.tail = self._next(self.tail)
28             self.size += 1
29             self.condition.notify()
30
31     def get(self):
32         with self.condition:
33             while self.is_empty():
34                 self.condition.wait()
35             item = self.buffer[self.head]
36             self.buffer[self.head] = None
37             self.head = self._next(self.head)
38             self.size -= 1
39             self.condition.notify()
40             return item
```

CAN YOU GUESS TYPE OF BUG?



QUEUE TLA SPEC: PART 1

```
1  ----- MODULE buffer -----
2  EXTENDS Naturals, Sequences
3
4  CONSTANTS Producers,
5             Consumers,
6             BufCapacity,
7             Data
8
9  ASSUME /\ Producers # {}
10         /\ Consumers # {}
11         /\ Producers \intersect Consumers = {}
12         /\ BufCapacity > 0
13         /\ Data # {}
14  VARIABLES buffer,
15             waitSet
16  -----
```

QUEUE TLA SPEC: PART 2

```
16 -----
17 Participants == Producers \union Consumers
18 RunningThreads == Participants \ waitSet
19
20 TypeInv == /\ buffer \in Seq(Data)
21             /\ Len(buffer) \in 0..BufCapacity
22             /\ waitSet \subseq Participants
23
24 Notify == IF waitSet # {}
25           THEN \E x \in waitSet : waitSet' = waitSet \ {x}
26           ELSE UNCHANGED waitSet
27
28 NotifyAll == waitSet' = {}
29
30 Wait(t) == waitSet' = waitSet \union {t}
31 -----
```

QUEUE TLA SPEC: PART 3

```
32  Init == buffer = <<>> /\ waitSet = {}
33  Put(t,m) == IF Len(buffer) < BufCapacity
34                THEN /\ buffer' = Append(buffer, m)
35                      /\ Notify
36                ELSE /\ Wait(t)
37                      /\ UNCHANGED buffer
38  Get(t) == IF Len(buffer) > 0
39                THEN /\ buffer' = Tail(buffer)
40                      /\ Notify
41                ELSE /\ Wait(t)
42                      /\ UNCHANGED buffer
```

QUEUE TLA SPEC: PART 4

```
43  Next == \E t \in RunningThreads : \/ t \in Producers /\ \E m \in Data : Put(t,m
44                                     \/ t \in Consumers /\ Get(t)
45
46  Prog == Init /\ [] [Next]_<<buffer, waitSet>>
47  -----
48  NoDeadlock == [] (RunningThreads # {})
49
50  THEOREM Prog => [] TypeInv /\ NoDeadlock
```

Invariant definition

QUEUE TRACE


TLC shows all steps that leads to *invariant violation*.

Name	Value
waitSet	{"p1", "p2", "p3", "c3", "c4", "c5"}
▼ ▲ <Action line 46, col 9 to line 47, col 62 of module b...	State (num = 26)
buffer	<<"m1", "m1">>
waitSet	{"p1", "p2", "p3", "c4", "c5"}
▼ ▲ <Action line 46, col 9 to line 47, col 62 of module b...	State (num = 27)
buffer	<<"m1">>
waitSet	{"p1", "p2", "p3", "c5"}
▼ ▲ <Action line 46, col 9 to line 47, col 62 of module b...	State (num = 28)
buffer	<< >>
waitSet	{"p1", "p2", "p3"}
▼ ▲ <Action line 46, col 9 to line 47, col 62 of module b...	State (num = 29)
buffer	<< >>
waitSet	{"p1", "p2", "p3", "c1"}
▼ ▲ <Action line 46, col 9 to line 47, col 62 of module b...	State (num = 30)
buffer	<< >>
waitSet	{"p1", "p2", "p3", "c1", "c2"}
▼ ▲ <Action line 46, col 9 to line 47, col 62 of module b...	State (num = 31)
buffer	<< >>
waitSet	{"p1", "p2", "p3", "c1", "c2", "c3"}
▼ ▲ <Action line 46, col 9 to line 47, col 62 of module b...	State (num = 32)
buffer	<< >>
waitSet	{"p1", "p2", "p3", "c1", "c2", "c3", "c4"}
▼ ▲ <Action line 46, col 9 to line 47, col 62 of module b...	State (num = 33)
buffer	<< >>
waitSet	{"p1", "p2", "p3", "c1", "c2", "c3", "c4", "c5"}

THE BUG

- Condition variable shared between producers and consumer
- On producer signal other producer make wake up instead of consumer
- Bug is very hard to reproduce since 30 specific steps need to happen

NOTIFY_ALL() VS NOTIFY()



If ..., you **share a condition variable** between multiple predicates, you **must always broadcast**, never signal;
Book: Programming with POSIX Threads by David R. Butenhof

MULTITHREADING QUEUE SPEC RESULTS

- Thread programming is hard.
- More than **30 steps** required to reproduce deadlock!
- Notify all strategy fixes issue with deadlock.
- If number of threads $> 2 * \text{buffer capacity}$ algorithm is not deadlock free.

AIORWLOCK SPEC

AIORWLOCK – READ WRITE LOCK FOR ASYNCIO

An RW lock allows concurrent access for read-only operations, while write operations require exclusive access, simple example:

```
1  import asyncio
2  import aiowlock
3
4  async def go():
5      rwlock = aiowlock.RWLock()
6
7      async with rwlock.writer:
8          print("inside writer: only one writer is possible")
9
10     async with rwlock.reader:
11         print("inside reader: multiple reader possible")
12
13 loop = asyncio.get_event_loop()
14 loop.run_until_complete(go())
```

AIORWLOCK – BUG

Failing tests for some corner-case uses #37

 Closed popravich wants to merge 1 commit into master from corner_case_tests

 Conversation 4

 Commits 1

 Files changed 1



popravich commented on May 19, 2017

Member



Several tests to show when `RWLock` doesn't work as expected.



 add failing tests for some corner-case uses

 b00c109



popravich commented on May 19, 2017


Member



@jettify, please take a look.

AIORWLOCK POSSIBLE STATES


Read Write lock state machine:



AIORWLOCK TRACE

Name	Value
▼ ▲ <Initial predicate>	State (num = 1)
▶ Lock	[t1 -> "Waiting", t2 -> "Waiting", t3 -> "Waiting"]
State	0
▼ ▲ <Action line 30, col 15 to line 33, col 40 of mod...	State (num = 2)
▼ Lock	[t1 -> "Write", t2 -> "Waiting", t3 -> "Waiting"]
● t1	"Write"
● t2	"Waiting"
● t3	"Waiting"
State	-1
▼ ▲ <Action line 26, col 18 to line 28, col 38 of mod...	State (num = 3)
▼ Lock	[t1 -> "WriteRead", t2 -> "Waiting", t3 -> "Waiting"]
● t1	"WriteRead"
● t2	"Waiting"
● t3	"Waiting"
State	0
▼ ▲ <Action line 22, col 18 to line 25, col 43 of mod...	State (num = 4)
▼ Lock	[t1 -> "WriteRead", t2 -> "Read", t3 -> "Waiting"]
● t1	"WriteRead"
● t2	"Read"
● t3	"Waiting"
State	1

AIORWLOCK POSSIBLE STATES FOR 2 TASKS



AIORWLOCK SPEC RESULTS

- In fact bug reveled itself on specification phase, before running any models.
- Only three steps required to reproduce issue.
- First aio-lib's project with formal verification!

CONCLUSIONS

LIMITATIONS OF MODEL-CHECKING

- State space explosion number of states reachable by a system can quickly become huge, or even infinite
- Used as an adjunct to, not a replacement for, standard quality assurance methods
- Formal methods are not a panacea, but can increase confidence in a product's reliability if applied with care and skill
- Very useful for consistency checks, but can not assure completeness

Questions?



<http://github.com/jettify>

REFERENCES i



azurewebsites.org.

Principles of chaos engineering, 2018.

<http://principlesofchaos.org>.



W. contributors.

Therac-25 — wikipedia, the free encyclopedia, 2018.

<https://en.wikipedia.org/w/index.php?title=Therac-25&oldid=820552180>.



Elasticsearch.

Formal models of core elasticsearch algorithms, 2017.

<https://github.com/elastic/elasticsearch-formal-models>.



A. Hadoop.

Introduce coordination engine interface, 2017.

<https://issues.apache.org/jira/browse/HADOOP-10641>.



lamport.

The tla toolbox home page, 2018.

<http://lamport.azurewebsites.net/tla/toolbox.html>.



lkml.org.

kernel/locking: qspinlock improvements, 2018.

<https://lkml.org/lkml/2018/4/5/496>.

REFERENCES ii



Microsoft.

A technical overview of azure cosmos db, 2017.

<https://azure.microsoft.com/en-us/blog/a-technical-overview-of-azure-cosmos-db/>.



mit.edu.

Mit 16.35 aerospace software engineering, 2002.

<http://web.mit.edu/16.35/www/lecturenotes/FormalMethods.pdf>.



C. Newcombe.

Why amazon chose tla+.

In Y. Ait Ameur and K.-D. Schewe, editors, *Abstract State Machines, Alloy, B, TLA, VDM, and Z*, pages 25–39, Berlin, Heidelberg, 2014. Springer Berlin Heidelberg.



Rapitasystems.

What really happened to the software on the mars pathfinder spacecraft?, 2013.

<https://www.rapitasystems.com/blog/what-really-happened-to-the-software-on-the-mars-pathfinder-spacecraft>.



I. Stoica, R. Morris, D. Karger, M. F. Kaashoek, and H. Balakrishnan.

Chord: A scalable peer to peer lookup service for internet applications.

ACM SIGCOMM Computer Communication Review, 31(4):149–160, 2001.



P. Zave.

How to make chord correct (using a stable base).

CoRR, abs/1502.06461, 2015.

REFERENCES iii



S. Zhou.

Tla+ spec of a simplified part of mongodb replication system, 2017.

<https://github.com/visualzhou/mongo-repl-tla>.